

Artem Sidnev

AI Engineer (Middle) | LLM Applications, RAG, Agentic Systems

Moscow • remote/hybrid • open to relocation

+7 910 969 29 22 | a.sidnevart@gmail.com | [Telegram](#) | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

Summary

Applied AI engineer with a strong production-backend foundation, focused on shipping LLM-powered products end to end. Built AI developer tooling (code review, CI refactoring, RAG onboarding) in a banking environment, an autonomous agentic knowledge system, and LLM features in open source. Comfortable owning the full path: prompt and retrieval design, integration, evaluation, deployment, and cost/latency tuning. Stack: Python, OpenAI, LangChain, pgvector, FastAPI, Docker.

Core Skills

LLM & GenAI: OpenAI API, prompt engineering, structured outputs, tool/function calling, human-in-the-loop, output evaluation

RAG & Retrieval: LangChain, pgvector, embeddings, semantic search, document retrieval, source citation, confidence scoring

Agentic Systems: multi-agent orchestration, MCP servers/tools, autonomous & scheduled agents, workflow automation (n8n)

AI Engineering / Ops: Docker, GitLab CI / GitHub Actions, observability (Grafana, Prometheus), token & cost optimization

Backend for AI: Python, FastAPI, Go, Java/Kotlin, PostgreSQL, Redis, Kafka, gRPC, REST

Languages: Python (primary), Go, TypeScript, Java/Kotlin, SQL

Experience

Software Engineer (Backend & AI Tooling) | T-Bank

Oct 2025 – Present

- Built an **AI-assisted code-review bot** for GitLab: an MR-analysis pipeline that summarizes changes, classifies issues by severity, and suggests fixes, with human-in-the-loop approval (OpenAI, GitLab API, FastAPI).
- Built an **AI CI refactor bot**: detects tech-debt patterns, generates refactoring plans, and prepares draft MRs in CI — accelerated the main pipeline by **5–7 min (15–20%)** and reduced flaky tests.
- Built a **RAG onboarding agent** over internal documentation: retrieval pipeline, context-aware chat, source citation, and confidence scoring (LangChain, PostgreSQL, Redis, FastAPI).
- Automated campaign workflows with **n8n** (creation, partner-integration monitoring, result reporting), saving **40–60 manual hours/month**.
- Backend depth behind the AI work: made audience gathering size-independent (**1–10 s**) and cut export from **8 h to 15 min (>32×)** on a **~30M-transactions/day** platform — I build the production systems that AI features run on.

Creator & AI Engineer | AI Office (OpenClaw), open-source

2025 – Present

- Built an autonomous agentic “second brain”: scheduled agents ingest blogs, YouTube, and podcasts, generate daily/weekly digests, draft content, and monitor career opportunities, writing into a structured knowledge base.
- Implemented **RAG over the knowledge base**, persistent memory via **MCP servers**, web-search augmentation, and Grafana cost/usage monitoring.
- Cut LLM token cost **~60%** (measured on own setup) through output compression and prompt/config optimization.
- Stack: OpenClaw, OpenAI, MCP, WebSocket, Grafana, Docker.

Full-stack & Applied AI | ProofForge, open-source

2024 – Present

- Built **AI-assisted goal refinement** and **AI-generated growth dossiers** with OpenAI, plus **semantic search** over user data via pgvector.
- Go backend (chi, pgx) and Next.js frontend; PostgreSQL/pgvector, Redis, MinIO, Docker, GitHub Actions CI/CD.

Backend & Automation Engineer | Client Projects

May 2022 – Aug 2025

- ~3 years building Python/FastAPI backends, Telegram bots and Mini Apps, and CRM/API integrations (real-estate analytics, booking & loyalty marketplaces) — the engineering groundwork behind the AI work above.

Education

Central University | B.Sc. in Computer Science | 2026 (*studying alongside work*)