Artem Sidnev

Backend Engineer | Java/Kotlin/Python, Spring Boot, PostgreSQL, ClickHouse, Kafka

Email | Telegram | LinkedIn | GitHub | Portfolio

Summary

Backend engineer with 3+ years of experience in production services, internal platforms, data-heavy workflows, and automation. Strongest in backend/platform work with complex business rules, queues, integrations, and large data volumes. I usually take problems from requirements and API design to implementation, rollout, metrics, and production support. Main stack: Java/Kotlin/Python, Spring Boot, PostgreSQL, ClickHouse, Kafka, REST/gRPC APIs, task queues, performance work, and pragmatic system design.

Core Skills

Languages: Java, Kotlin (JVM), Python, Go, TypeScript, SQL
Backend: Spring Boot, REST APIs, gRPC, microservices, API design, unit and integration testing, Agile, code review, domain modeling, complex business logic, validation, error handling
System Design: distributed systems, service boundaries, REST/gRPC contracts, async pipelines, idempotency, task states and cancellation, bounded concurrency, backpressure, feature flags, fallback paths
Data & Messaging: PostgreSQL, ClickHouse/CHaaS, Redis, Kafka, Postgres-backed task queues, S3, SQL diagnostics
Infrastructure: Docker, Docker Compose, Kubernetes, nginx, GitLab CI/CD, monitoring and observability (Prometheus, OpenTelemetry, Grafana)
Product Domains: targeting platforms, partner integrations, analytics systems, booking/loyalty workflows, AI automation, RBAC/role-based flows

Experience

Software Engineer (Backend) | T-Bank *Oct 2025 – Present*

- Build backend systems for a manager-facing cashback audience platform that selects customers for partner-funded offers using transaction, customer, and partner data.
- Designed a control-plane/data-plane split for audience gathering: the Kotlin/Spring service kept business orchestration, validation, and lifecycle ownership, while heavy materialization/export workloads were delegated to an internal Go execution backend via gRPC, ClickHouse, and S3.
- Built a contract-first async integration with external task lifecycle, typed error mapping, idempotency boundaries, cancellation/retry semantics, feature-flagged rollout, legacy fallback, and task-level observability.
- Identified and adapted a reusable internal platform capability from another team, applying bucketized ClickHouse processing, CH-to-S3 streaming, task-type isolation, and bounded concurrency to the audience gathering pipeline.
- Built a quantitative latency model for the Postgres-backed DAG pipeline, decomposing total time into queue wait, execution, scheduler polling gaps, and critical-path depth.
- Reworked DAG structure, queue isolation, concurrency limits, and ClickHouse workload management: audience gathering became size-independent, completing in 1–10 seconds for any audience (down from 6+ minutes typical, hours for large ones), from initiative to production rollout.
- Cut Auditory Export time from 8+ hours to 15 minutes (>32x), eliminating the bottleneck between audience gathering and partner platform delivery.
- Made campaign launches less dependent on heavy audience-building jobs: faster audience refresh, fewer technical stalls, and more accurate inclusion/exclusion of customers in cashback offer audiences.
- Delivered spend/income targeting on about 30M transactions per day and 1.5 years of historical data; estimated product impact is up to RUB 7.5M per year.
- Moved audience-building workflows to an internal ClickHouse-as-a-Service environment, making a critical scenario about 5x faster and reducing dependence on the legacy stack.
- Automated partner-region substitution and early validation in targeting flows, saving up to 40–60 manual hours per month and cutting common issue diagnosis by 20–40 minutes.
- Improved engineering productivity by accelerating the main CI pipeline by 5–7 minutes (15–20%), reducing flaky tests, and building AI-assisted tools for GitLab review, CI refactoring, n8n automation, and RAG onboarding.
- Fixed a production document-generation issue that produced incorrect data in documents, removing a source of manual support follow-ups and repeated fixes.

Backend and Automation Engineer | Freelance / Client Projects *May 2022 – Present*

- Built a commercial real estate analytics platform combining auction data, CIAN listings, geocoding, a Telegram bot, an admin panel, and a searchable property catalog.
- Reduced property economics evaluation to about 10 minutes and district-level market analysis to 7–10 minutes by automating data collection, normalization, comparables, yield, cash flow, and payback calculations.
- Delivered booking, loyalty, and service-slot workflows with REST APIs, PostgreSQL, Redis, Telegram Mini Apps, Docker/nginx deployment, YClients integration, webhooks, retries, and deployment automation.
- Built customer, partner, and admin flows for service businesses; launched an MVP with 3 large partners, reducing slot vacancy by 35%, increasing partner revenue by 20%, and improving booking conversion by 25%.
- Implemented matching, Elo rating, bonus debit/credit settlement, contact-based registration, and role-aware operational flows for partner/customer products.

Education

Central University | B.Sc. in Computer Science, in progress
Additional: Harvard CS50, MIT OCW Algorithms, freeCodeCamp Java/Spring